

Our solution to the IDAO 2020 qualifiers

Max Halford
Raphaël Sourty
Robin Vaysse

University of Toulouse
Sunday 12th April, 2020



Our team


- Max Halford, 3rd year PhD student at IMT/IRIT
- Raphaël Sourty, 1st year PhD student at IRIT
- Robin Vaysse, 1st year PhD student at IRIT

We like competitive data science!

Context

- Satellite position forecasting
- Two tracks with separate leaderboards:
 1. Make the most accurate predictions possible
 2. Make accurate predictions with two constraints:
 - 2.1 Take less than 60 seconds
 - 2.2 Keep peak RAM usage under 500MB


The data



Our solution in a nutshell

- We train one model per satellite and per coordinate ($300 \times 6 = 1800$ models)
- Each model is an autoregressive (AR) process of order $p = 48$
- In other words, we train a linear regression to predict y_{n+1} from $\{y_{n-48}, \dots, y_n\}$, that's all!
- To predict several steps ahead, we use the prediction at step $n + 1$ as a feature at step $n + 2$
- We validate locally on the last 40% of the data
- Our approach is simple enough to be used for both tracks without modifications

Starting simple



Auto-regression

- Using past target values makes sense because the data is very periodic
- For every satellite and coordinate, we build a vector of features
- Each vector contains the p past target values
- We obtain n feature vectors and n targets
- For forecasting into the future, we:
 1. Make a prediction for the next time step
 2. Append the prediction to the feature vector
 3. Remove the oldest value from the vector
 4. Repeat from step 1.
- Flexible framework:
 - Any regression model can be plugged in
 - Any feature can be added, provided it can be computed online

Dealing with speed

- AR models are slow at inference because of their sequential nature
- In `scikit-learn`, calling `.predict(X)` many times incurs a large overhead
- We “stripped” the `scikit-learn` classes we used to their bare minimum by overriding some of their methods

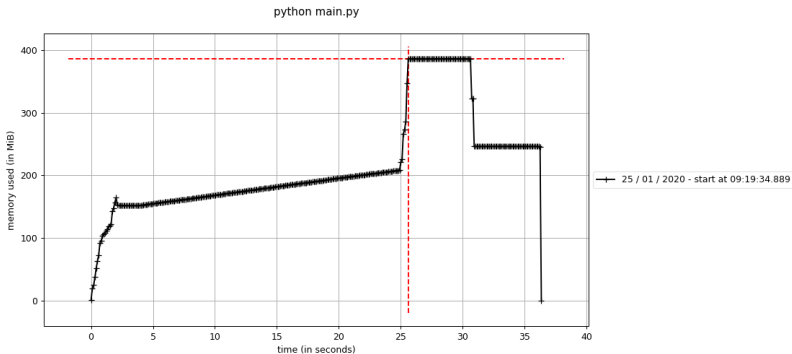
Overriding scikit-learn's linear regression

```
class StandardScaler(preprocessing.StandardScaler):  
    """Barebones implementation with less overhead than sklearn."""  
  
    def transform(self, X):  
        return (X - self.mean_) / self.var_ ** .5  
  
class LinearRegression(linear_model.LinearRegression):  
    """Barebones implementation with less overhead than sklearn."""  
  
    def predict(self, X):  
        return np.dot(X, self.coef_) + self.intercept_
```

More information [here](#). We've also learned about [sklearn-onnx](#).

Dealing with memory usage

We used a Python package called `memory_profiler` to measure the memory usage of our script.



What didn't work

- Gaussian processes with sinusoidal kernels gave good training results, but fared poorly on the test set
- The N-BEATS¹ model fits perfectly to the training data but diverges in auto-regressive mode
- We got no improvement by training a multi-output linear regression to try capturing coordinate dependencies

¹Boris N. Oreshkin et al. “N-BEATS: Neural basis expansion analysis for interpretable time series forecasting”. In: *CoRR* abs/1905.10437 (2019). arXiv: [1905.10437](https://arxiv.org/abs/1905.10437). URL: <http://arxiv.org/abs/1905.10437>.

Production considerations

- Our model is essentially a linear regression
- Linear regression can be trained with stochastic gradient descent (SGD)
- SGD requires one sample at a time, and is thus enables online algorithm
- Online learning allows learning from a stream of data
- Predicting satellite positions is inherently a streaming problem, therefore models that can be trained online should be preferred

Shameless publicity: check out [creme](#) and [chantilly](#) for online learning

Our advice for competitive data science

- *“Keep it simple, stupid”* (KISS principle)
- Always start by setting up a local validation benchmark
- When your model improves, save your work (`git` is your friend)
- Doubt everything you do
- Don't be scared to try stuff, but don't tunnel vision

Code can be found [on GitHub](#)

Thank you for listening!